



redhat.

# Incremental Backups

*(Good things come in small packages!)*

John Snow (*yes, I know*)  
Software Engineer, Red Hat  
2017-02-05

# Acknowledgments

(Because computers are awful and I need help sometimes)

No feature is an island, so I'd like to acknowledge:

- Jagane Sundar
  - Initial feature proposal and prior work (2011)
- Fam Zheng
  - Initial drafts for current version (2014-2015)
- Stefan Hajnoczi & Max Reitz
  - Reviews and patience

# Acknowledgments

(Because computers are awful and I need help sometimes)

No feature is an island, so I'd like to acknowledge:

- Vladimir Sementsov-Ogievskiy, Virtuozzo
  - Advanced features (Persistence, Migration)
  - Performance enhancements
  - Reviews, Patience, and general excellence
- Denis Lunev, Virtuozzo
  - Dedicated and persistent involvement

# Overview

(Things I hope not to stammer through)

## Prologue

- Problem Statement
- Approach
- Design Goals

## Act I: Building Blocks

- Block Dirty Bitmaps
- QMP interface and usage
- QMP transactions

# Overview

(Things I hope not to stammer through)

## Act II: Life-cycle

- Incremental backup life-cycle
- Examples

## Aside: Transactions

- BlockJobs
- Transactions
- Multi-drive Coherency
- Errors

# Overview

(Things I hope not to stammer through)

## Act III: Advanced Features

- Migration
- Persistence
- Push vs Pull model backups
- TODOs

## Dénouement

- Project Status, Questions and Answers

# PROLOGUE

(In which our heroes come to know the enemy)

# The Problem

(I just wandered into this talk, what's it about?)



Monday  
128GiB



Tuesday  
128GiB



Wednesday  
128GiB



Thursday  
128GiB



Friday  
128GiB

**Gross.**

- Abysmal storage efficiency
- Clunky, slow
- But admittedly simple and convenient

# The Problem

(I just wandered into this talk, what's it about?)



Monday  
128GiB



Tuesday  
2GiB



Wednesday  
2.5GiB



Thursday  
2.21GiB



Friday  
1GiB

Much Better!

- Efficient: only copies modified data
- Fast!
- More complicated...?

# Welcome!

(You're in my world now)

QEMU added preliminary support for incremental backups in QEMU 2.4, 2015-08-11.

- (I can't commit to either US or EU dates, so enjoy this ISO one instead)
- Development is ongoing as of 2.8
- Not included as “supported” in a Red Hat product yet
  - So, it's mostly for the brave.
  - But we're nearing feature completion.

# Approach

(Where did we come from; where did we go)

Incremental Live Backups have a storied lineage.

- Jagane Sundar's LiveBackup (2011)
  - Separate CLI tools
  - Entirely new network protocol
  - Ran as an independent thread
  - Utilized temporary snapshots for atomicity
  - Implemented with in-memory dirty block bitmaps
  - Was ultimately not merged

# Approach

(Where did we come from; where did we go)

## Fam Zheng's Incremental Backup (2014)

- Also dirty sector bitmap based
  - Uses existing HBitmap/BdrvDirtyBitmap primitives
- No new external tooling or protocols
- Managed via QMP
- Implemented simply as a new backup mode
- Can be used with any image format
- Maximizes compatibility with existing backup tools

# Design Goals

(What do we want?)

- Reuse existing primitives as much as possible
  - Key structure: 'block driver dirty bitmap'
    - Already tracks dirty sectors
    - Used for drive mirroring, block migration
    - Configurable granularity
    - Many bitmaps can be used per-drive

# Design Goals

(What do we want? Efficient Backups!)

- Reuse existing primitives
  - Key interface: drive-backup
    - Implemented via well-known QMP protocol
    - Used to create e.g. full backups
    - Already capable of point-in-time live backups
    - Can already export data via NBD
    - We merely add a new sync=incremental mode
      - ...And a bitmap=<name> argument.

# Design Goals

(When do we want it?)

- Coherency
  - Multi-drive point-in-time backup accuracy
  - Utilize existing QMP transaction feature
- Persistence
  - Bitmaps must survive shutdowns and reboots
  - Must not depend on drive data format
  - Nor on the backup target format

# Design Goals

(When do we want it? By 2.9 hopefully!)

- Migration-safe
  - Migrating must not reset or lose bitmap data
- Error Handling
  - Bitmap data must not be lost on backup failure
  - Starting a new full backup is not sufficiently robust
- Integrity
  - We *must* be able to detect desync between persistence data and block data

# Why not use snapshots?

(Saving you time during the Q&A)

“Both offer point-in-time views of data, why not use the existing mechanism?”

- No need to parse format-specific snapshots on disk
- We can use *any format*
- Incremental backups are *inert* and do not grow
  - No IO required to delete incrementals
- We can utilize existing backup frameworks
- Access to QEMU's NBD server

# ACT I: BUILDING BLOCKS

(In which our heroes prepare for battle)

# Block Dirty Bitmaps

(Nothing to do with your image search settings)


Before showcasing incrementals, some background:

- BdrvDirtyBitmap is the existing block layer structure used to track writes
  - Already used for drive-mirror, live block migration
  - Implemented using hierarchical bitmap
  - Any number can be attached to a drive
    - Allows for multiple independent backup regimes

# Block Dirty Bitmaps

(Nothing to do with your image search settings)


Hbitmap hierarchy:



# Block Dirty Bitmaps

(Nothing to do with your image search settings)

Bitmap plurality:



# Block Dirty Bitmaps - Naming


(A bitmap by any other name would smell as sweet...?)

- Block dirty bitmaps may have names:
  - Existing internal usages are anonymous
  - The name is unique to the drive
  - Bitmaps on different drives can have the same name
  - The (node, name) pair is the bitmap ID
    - Used to issue bitmap management commands

# Block Dirty Bitmaps - Naming

(A bitmap by any other name would smell as sweet...?)

Bitmap naming:



# Block Dirty Bitmaps - Granularity

(Backups from *French Press* to *Turkish*)

- Block dirty bitmaps have granularities:
  - Small granularity – smaller backups\*
  - Uses more memory
    - 1 TiB w/ g=32KiB → 4MiB
    - 1 TiB w/ g=128KiB → 1MiB
  - Default: 64KiB\*\*
    - Attempts to match cluster size
    - 64KiB clusters (default) for qcow2

# Granularities – In Detail

(Tuned like the finest \$4 ukulele)

- Bitmaps track writes *per-sector*
  - Configure granularity in **bytes**
  - 64K → 128 sectors (512 bytes/sector)
- The backup engine itself copies out per-cluster
  - Currently: non-configurable, 64K clusters
- The file format also has a cluster size
  - qcow2 defaults to 64K.
- Conclusion: 64K is probably best (for now)

# Block Dirty Bitmaps - Management

(Bitmap wrangling 101)

We need to manage these bitmaps to make backups.

- Managed via QMP
  - Good news if you're a computer!
  - Four commands:
    - `block-dirty-bitmap-add`
    - `block-dirty-bitmap-remove`
    - `block-dirty-bitmap-clear`
    - `query-block`

# Block Dirty Bitmaps - Creation

(Let there be... bits!)

- Bitmaps can be created at any time, on any node
- Bitmaps begin recording writes immediately
- Granularity is optional

```
{ "execute": "block-dirty-bitmap-add",
  "arguments": {
    "node": "drive0",
    "name": "bitmap0",
    "granularity": 131072
  }
}
```

# Block Dirty Bitmaps - Deletion

(For days when *less* is *more*)

- Can only be deleted when not in use
- Bitmaps are addressed by their (node, name) pair
- Has no effect on backups already made
- Has no effect on other bitmaps or nodes

```
{ "execute": "block-dirty-bitmap-remove",
  "arguments": {
    "node": "drive0",
    "name": "bitmap0"
  }
}
```

# Block Dirty Bitmaps - Resetting

(Sometimes we just want a second chance)

- Bitmaps can be cleared of all data
- Primarily for convenience
- Begins recording new writes immediately, like add

```
{ "execute": "block-dirty-bitmap-clear",
  "arguments": {
    "node": "drive0",
    "name": "bitmap0"
  }
}
```

# Block Dirty Bitmaps - Querying

(Who are you? Who who, who who?)

Bitmap data can be retrieved via block-query.

```
{"execute": "query-block", "arguments": {}}

{"return": [{ ...
  "device": "drive0",
  "dirty-bitmaps": [{{
    "status": "active",
    "count": 296704,
    "name": "bitmap0",
    "granularity": 65536 }]
  ... }]}
```

# Block Dirty Bitmaps - Querying

(Who are you? Who who, who who?)

Bitmap data can be retrieved via block-query.

```
{"execute": "query-block", "arguments": {}}

{"return": [{ ...
  "device": "drive0",
  "dirty-bitmaps": [{ ...
    "status": "active",           (or "frozen"!)
    "count": 296704,
    "name": "bitmap0",
    "granularity": 65536 }]
... }]}
```

# Block Dirty Bitmaps - Querying

(Who are you? Who who, who who?)

Bitmap data can be retrieved via block-query.

```
{"execute": "query-block", "arguments": {}}

{"return": [{ ...
  "device": "drive0",
  "dirty-bitmaps": [{ ...
    "status": "active",
    "count": 296704,          (sectors!)
    "name": "bitmap0",
    "granularity": 65536 }]} (2318 clusters)
... }]}
```

# Building Cognitive Dissonance

(Problem Statement 2: Electric Boogaloo)

- QMP commands are not particularly useful alone
  - They are not atomic
  - Only “safe” when VM is offline
  - No cross-drive coherence guarantee

# Incremental Transactions

(Dissonance abated!)

- Bitmap management transactions allow us to–
  - Create full backups alongside a bitmap reset
  - Create a full backup alongside a new bitmap
  - Reset bitmaps across multiple drives
  - Issue a number of incremental backups across multiple drives

# Incremental Transactions


(Dissonance abated!)

- Supported transaction actions:
  - type:block-dirty-bitmap-add
  - type:block-dirty-bitmap-clear
- No transaction needed for remove
- Works in conjunction with type:drive-backup
  - For incrementals (multi-drive coherency)
  - For full backups
    - new incremental chains / sync points

# ACT II: LIFE CYCLE

(In which our heroes save time and money)

# Incrementals – Life Cycle



- 1) Create a new backup chain, or
- 2) Synchronize an existing backup chain
- 3) Create the first incremental backup
- 4) Create subsequent incremental backups

# Life Cycle – New Chain

(There and backup again)

Example 1: Start a new backup chain atomically

```
{ "execute": "transaction",
  "arguments": {
    "actions": [
      {"type": "block-dirty-bitmap-add",
       "data": {"node": "drive0", "name": "bitmap0"} },
      {"type": "drive-backup",
       "data": {"device": "drive0",
                 "target": "/path/to/full.qcow2",
                 "sync": "full", "format": "qcow2"} }
    ]
  }
}
```


# Life Cycle – New Chain

(There and backup again)

**id=drive0**

# Life Cycle – New Chain

(There and backup again)



# Life Cycle – New Sync Point


(Sunday night maintenance blues)

Example 2: Take an existing bitmap and create a new full backup as a synchronization point.

```
{ "execute": "transaction",
  "arguments": {
    "actions": [
      {"type": "block-dirty-bitmap-clear",
       "data": {"node": "drive0", "name": "bitmap0"} },
      {"type": "drive-backup",
       "data": {"device": "drive0",
                 "target": "/path/to/new_full_backup.qcow2",
                 "sync": "full", "format": "qcow2"} }
    ]
  }
}
```

# Life Cycle – New Sync Point


(Sunday night maintenance blues)



**full.qcow2**

# Life Cycle – New Sync Point

(Sunday night maintenance blues)



# Life Cycle – First Incremental

(The first step of our journey)


Example 3: Create an incremental backup. Can be done via transaction or single QMP command.

```
# qemu-img create -f qcow2 inc.0.qcow2 -b full.qcow2 -F qcow2
```

```
{ "execute": "drive-backup",
  "arguments": {
    "device": "drive0",
    "bitmap": "bitmap0",
    "target": "inc.0.qcow2",
    "format": "qcow2",
    "sync": "incremental",
    "mode": "existing"
  }
}
```


# Life Cycle – First Incremental

(The first step of our journey)



# Life Cycle – First Incremental

(The first step of our journey)



# Life Cycle – Subsequent Backups

(To infinity, and beyond!)


Examples [4, $\infty$ ): Create subsequent incrementals.

```
# qemu-img create -f qcow2 inc.<n>.qcow2 -b inc.<n-1>.qcow2 -F qcow2
```

```
{ "execute": "drive-backup",
  "arguments": {
    "device": "drive0",
    "bitmap": "bitmap0",
    "target": "inc.<n>.qcow2",
    "format": "qcow2",
    "sync": "incremental",
    "mode": "existing"
  }
}
```


# Life Cycle – Subsequent Backups

(To infinity, and beyond!)



# Life Cycle – Subsequent Backups

(To infinity, and beyond!)



# Interlude

# Interlude: Transactions

(Just kidding, we're gonna talk about more stuff)

# Explainer: Block Jobs

(Jobs & The Economy: Redux)

- What are jobs? (ha ha ha)
  - QMP commands are synchronous
  - QMP socket blocks on each command
  - So what about long-running commands?
- BlockJobs: Asynchronous task API
  - Allows management via further QMP commands
  - For more info: See literally\* any talk from KVM Forum 2016

\*figuratively

# Transactions - detail

(In case you forgot? Sorry, there's a lot of stuff.)

## Transactions:

- Allow batching of certain QMP commands
- Each individual item is an “action”
- Transaction succeeds only if all actions do
- Some actions/commands launch jobs
- Some do not.
- Wow, I hope that doesn't cause any problems.

(Of course it did.)

# Transactions X Jobs

(Transaction Interaction Intersection)

How do job-actions work?

- Before 2.5:
  - Action succeeds if job is *started*
  - Jobs failing later have no effect on other jobs
  - Some backups succeed, some fail
  - `completion_mode=individual`

# Transactions X Jobs

(Transaction Interaction Intersection)

How do job-actions work?

- After 2.5, with completion\_mode=grouped ...
  - Action succeeds if job is started
    - No change from ‘individual’ mode
    - Jobs cannot complete until **all** jobs ready to
    - One job will cause all others to fail
  - Clients can avoid keeping state on partial failures


# Multidrive Coherency

(Transaction actions in action (not to be confused with inaction))



```
{ "execute": "transaction",
  "arguments": {
    "actions": [
      { "type": "drive-backup",
        "data": { "device": "drive0", "bitmap": "bitmap0",
                  "format": "qcow2", "mode": "existing",
                  "sync": "incremental",
                  "target": "inc0.a.qcow2" } },
      { "type": "drive-backup",
        "data": { "device": "drive1", "bitmap": "bitmap1",
                  "format": "qcow2", "mode": "existing",
                  "sync": "incremental",
                  "target": "incl.a.qcow2" } },
    ]
  }
}
```

# Multidrive Coherency

(Twice as nice!)




**full0.qcow2**



**full1.qcow2**


# Multidrive Coherency

(Thrice as nice?)




# Multidrive Coherency

(...frice?)





# Partial Failures, Individual

(Not *my* problem)



**full0.qcow2**

**inc0.a.qcow2**




**full1.qcow2**

**inc1.a.qcow2**


# Partial Failures, Individual

(Not *my* problem)




# Partial Failures, Grouped

(Stronger together?)



**full0.qcow2**

**inc0.a.qcow2**





**full1.qcow2**

**inc1.a.qcow2**



# Partial Failures, Grouped

(Stronger together?)



**full0.qcow2**

**inc0.a.qcow2**





**full1.qcow2**



**inc1.a.qcow2**

# Partial Failures, Grouped

(Stronger together?)



**full0.qcow2**



**full1.qcow2**

# ACT III: ADVANCED FEATURES

(In which our heroes *rise above*)

# Bitmap Migration - 1<sup>st</sup> attempt

(Pack your data, we're moving to <target>)

- Mechanism similar to disk migration
- Data split into chunks (1KiB)
  - Bitmaps serialized piece-by-piece
- For sets of bitmaps below 1MiB...
  - Skip the live phase and copy the data wholesale.
  - 64GiB disk bitmap is only 128KiB
    - (+node and bitmap names, and stream metadata)

# Bitmap Migration - 1<sup>st</sup> attempt

(Pack your data, we're moving to <target>)

- Bitmaps not transferred alongside data
  - Transferred separately for flexibility
  - “meta bitmaps” (*dirty “dirty bitmap” bitmaps!?*)
    - Captures changes during live migration
    - Pieces can be resent if needed.
    - Uses *very little* memory: 64GiB → 16 bytes

# Bitmap Migration - 2<sup>nd</sup> attempt

(We're on the road again...)

- 1<sup>st</sup> approach worsens convergence problem
  - May not scale well
- New approach uses a post-copy technique
  - Simply send the whole bitmap post-pivot
  - Record new writes on target
    - Prohibit backups until data arrives
    - Re-merge bitmaps on target

# Bitmap Migration - Failures

(Mission Failed! We'll get 'em next time.)

What happens if the source dies post-pivot?

- Considered non-critical loss
- Bitmap chains can be re-started
- Future:
  - Reconstruct bitmap from two images?
- Other Options:
  - Use shared-storage migration
    - With persistence <stay tuned>

# Bitmap Persistence – Change of Plans

(I have altered the code. Pray I do not alter it further.)

- Plans *were* for a format-agnostic format
  - Using qcow2 to store bitmaps for arbitrary files
  - Plans scrapped...
- Now, we're targeting qcow2
  - More on other formats in a bit...!

# Bitmap Persistence

(Object permanence: not just for toddlers)

- Persistence targets the qcow2 format.
  - Multiple bitmaps can be stored per-file
  - Bitmaps have ‘types,’ we use a ‘dirty’ bitmap
  - Bitmaps can ‘autoload’ in QEMU
  - Spec amendment is merged!
  - Patches ready on-list from Virtuozzo

# Bitmap Persistence – Non qcow2

(AKA, “Can I please use this with raw?”)

- We have some options for other formats.
- Some formats may add primary support
  - Virtuozzo has expressed interest for parallels
  - Qcow2 with write-forwarding backing files?
    - Instead of read-only
    - Offer to forward writes
    - Allow for any format
    - Other benefits

# “Push Model” backups

(Let's take all our problems... and push them somewhere else!)

Backups described so far are “Push” model:

- QEMU “pushes” the data to a target
- It knows what sectors need to be pushed
- This works out pretty OK, but...
  - Some vendors wanted a different model

# “Pull Model” backups

(sometimes it's nice when doors work both ways)

The “Pull model” is different:

- QEMU offers a temporary, lightweight snapshot
  - “Image Fleecing”
  - Exported via NBD
- Via NBD extensions, client queries for status
- Client controls data flow
- Snapshot is deleted on close

# “Pull Model” backups

(sometimes it's nice when doors work both ways)

- Snapshot view is point-in-time
  - (like push model)
- Requires on-disk cache
- Offers full control on what is copied
  - How the data is stored is decided by the client
  - Most “QEMU-agnostic” method
- Only way to query dirty blocks

# TODOs

(<TODO: insert cheeky joke>)

- QMP interface for “pull” model
- QMP interface for modifying persistence attributes
- CLI tools for verification, analysis
  - Deletion/cleaning tools
  - “Offline” incremental backup support?
- “fsck support”
  - qemu-img check -r (?)

# TODOs

(<TODO: insert cheeky joke>)

- Data integrity
  - Periodic/opportunistic flushing
- GSOC / Outreachy 2017:
  - Reference implementation
  - CLI backup tool
  - Python?
  - Keep your eyes peeled:
  - [http://wiki.qemu.org/Google\\_Summer\\_of\\_Code\\_2017](http://wiki.qemu.org/Google_Summer_of_Code_2017)

# Dénouement

(In which our heroes live incrementally ever after)

# Project Status

(When do we get to use it!?)

- block-dirty-bitmap QMP interface
- sync=incremental mode (*push*)
- Transactions
- Qcow2 Persistence (Spec)
- Grouped Transactions
- Migration
- Persistence
- Pull model
- Merged! (2.4)
- Merged! (2.4)
- Merged! (2.5)
- Merged! (2.6)
- Merged! (2.8)
- Review, (2.9)
- Review, (2.9)
- Specs, (2.10+)

A wide-angle photograph of a mountainous landscape featuring extensive terraced rice fields. The fields are arranged in numerous parallel, wavy rows that follow the contours of the hillside. A small, simple wooden hut is positioned in the middle ground, nestled among the fields. The background shows more of the mountain range under a clear sky.

# Questions?

# Further Reading:

QEMU project wiki:

[http://qemu-project.org/Main\\_Page](http://qemu-project.org/Main_Page)

Bitmaps Documentation:

<.../qemu/docs/bitmaps.md>

QEMU iotests:

<.../qemu/tests/qemu-iotests/124>

Project status whitepaper (PDF):

<http://goo.gl/tT6n8S>

KVM Forum 2016 ‘jobs’ talk:

[http://events.linuxfoundation.org/sites/events/files/slides/kvm2016\\_v16.pdf](http://events.linuxfoundation.org/sites/events/files/slides/kvm2016_v16.pdf)

# THANK YOU!

More questions?  
[jsnow@redhat.com](mailto:jsnow@redhat.com)  
cc: [qemu-devel@nongnu.org](mailto:qemu-devel@nongnu.org)